



Department of Information Engineering and Computer
Science

Bachelor's Degree in
Information, Communications and Electronics Engineering

ROBOTICS PROJECT

Lecturer

Prof. Luigi Palopoli

Prof. Niculae Sebe

Prof. Michele Focchi

Prof. Placido Falqueto

Students

Mattia Meneghin [210561]

Filippo Conti [218297]

Nicola Gianuzzi [209309]

Academic Year 2022/2023

Contents

1	Introduction	1
1.1	Project Proposal 1	1
1.1.1	Assignment 1	1
1.1.2	Assignment 2	1
1.1.3	Assignment 3	1
1.1.4	Assignment 4	1
1.2	Delivery Rules	2
1.3	GitHub Repository	2
2	Set up your Environment	3
3	Simulation Environment	4
3.1	ROS	4
3.2	Catkin	4
3.2.1	Source Space	5
3.2.2	Build Space	5
3.2.3	Development (Devel) Space	5
3.2.4	Install Space	5
3.2.5	Result space	5
3.3	Locosim	5
3.4	Gazebo	6
3.5	Rviz	7
4	Motion	8
4.1	Planner	8
4.1.1	Planner commands from vision	8
4.1.2	motion/src/planner.cpp	8
4.1.3	motion/msg/legoFound.msg	8
4.2	Movement	10
4.2.1	motion/src/movement.cpp	10
4.2.2	motion/msg/legoTask.msg	10
4.2.3	Movement commands from the planner module	11
4.3	Kinetics	11
4.4	Acknowledgement system	12
4.5	spawnLegos.cpp	13
5	Vision	14
5.1	Tool used	14
5.1.1	Yolov5	14
5.1.2	Pytorch	14
5.1.3	Custom Roboflow version of Google Colab	14
5.1.4	MakeSense AI	15

5.2	Steps	16
5.2.1	Create Dataset	16
5.2.2	Dataset structure	16
5.2.3	Create Labels for the annotations	18
5.2.4	Create annotations	18
5.2.5	Test and refine the dataset	18
5.2.6	Roboflow Dataset Settings	19
5.2.7	Train into Google Colab provided by Ultralytics	20
5.3	Results	21
5.4	Dataset	22
5.5	In practice	23
5.5.1	Instance of assignment 2	23
5.5.2	Normal Vision Results	24
5.5.3	ROI Vision Results	25
6	Run The Project	26

Abstract

This project is based on ROS, namely Robot Operating System. More in detail we utilized the noetic ROS version and Locosim, a didactic framework to learn/test basic controllers schemes on quadruped robots. The project is composed by two main modules: **Vision** and **Movement**.

As for the vision we used YOLOv5 and PyTorch framework for the real time object detection through the ZED-Camera. The simulation was done in Gazebo and Rviz, exploiting ROS functionalities. All of these operations are possible thanks to Catkin, a build system of ROS, that combines CMake macros and Python scripts to provide some functionality on top of CMake's normal workflow and it was necessary in order to create the ROS nodes.

1 Introduction

1.1 Project Proposal 1

A number of objects (e.g., mega-blocks) are stored without any specific order on a stand (initial stand) located within the workspace of a robotic manipulator. The manipulator is an anthropomorphic arm, with a spherical wrist and a three-fingered gripper as end-effector. The objects can belong to different classes but have a known geometry (coded in the STL files). The objective of the project is to use the manipulator to pick the objects in sequence and to position them on a different stand according to a specified order (final stand). A calibrated 3D sensor is used to locate the different objects and to detect their mutual position in the initial stand. The project is organised as a sequence of assignments of increasing complexity.

1.1.1 Assignment 1

There is only one object in the initial stand, which is positioned with its base “naturally” in contact with the ground. The object can be of any of the classes specified by the project. Each class has an assigned position on the final stand, which is marked by a coloured shape representing the silhouette of the object.

KPI 1-1 time to detect the position of the object

KPI 1-2 time to move the object between its initial and its final positions, counting from the instant in which both of them have been identified.

1.1.2 Assignment 2

There are multiple objects on the initial stand, one for each class. There is no specific order in the initial configuration, except that the base of the object is “naturally” in contact with the ground. Each object has to be picked up and stored in the position prescribed for its class and marked by the object’s silhouette.

KPI 2-1 Total time to move all the objects from their initial to their final positions.

1.1.3 Assignment 3

There are multiple objects on the initial stand, and there can be more than one object for each class. The objects are positioned randomly on the stand but would not stand or lean on each other. An object could be lying on one of its lateral sides or on its top. Each object has to be stored in the position prescribed by its class. Objects of the same class have to be stacked up to form a tower.

KPI 3-1 Total time to move all the objects from their initial to their final positions.

1.1.4 Assignment 4

The objects on the initial stand are those needed to create a composite object with a known design (e.g., a castle). The objects are positioned randomly on the stand. An object could be lying on one of its lateral sides or on its top. The objects could also stand or lean on each other. The manipulator has to pick them up in sequence and create the desired composite object on the final stand.

1.2 Delivery Rules

The project is developed in groups. The typical group size consists of three-four members. We can also accept groups with a smaller number of members. The group is supposed to work in perfect cooperation and the workload is required to be fairly distributed. The specific contribution of each member will be exposed during the project discussion. The delivery phase is as follows:

1. The project will have to be tested in the laboratory with the Teaching Assistant at most 5 least five days before the exam date. During the tests, small videos can be shot and used for the presentation.
2. Each group will have to deliver the package containing the full code (with doxygen documentation and a readme for use) plus a 5-6 pages report describing
 - (a) the technique used for perception
 - (b) the technique used for robot motion
 - (c) the technique used for high-level planning
 - (d) A table with the KPI measured on Gazebo
3. The delivery deadline is three days before the exam presentation
4. On the day of the exam, the students will give a 20 minutes presentation highlighting the contribution of each member. A discussion will follow in which all members are supposed to answer questions on the entire project (regardless of her/his specific assignment within the group).
5. If allowed by the time, the group could also be asked to perform a small demo session. Otherwise, we will rely on the clip shot before the exam.

1.3 GitHub Repository

[Robotics Repository](#)

2 Set up your Environment

3 Simulation Environment

3.1 ROS

ROS is an open-source, meta-operating system for your robot. It provides the services you would expect from an operating system, including hardware abstraction, low-level device control, implementation of commonly-used functionality, message-passing between processes, and package management. It also provides tools and libraries for obtaining, building, writing, and running code across multiple computers. ROS is similar in some respects to "robot frameworks". [reference](#)

3.2 Catkin

Catkin packages can be built as a standalone project, in the same way that normal cmake projects can be built, but catkin also provides the concept of workspaces, where you can build multiple, interdependent packages together all at once. A catkin workspace is a folder where you modify, build, and install catkin packages. [reference](#)

```
workspace_folder/
  src/
    CMakeLists.txt      -- WORKSPACE
    package_1/
      CMakeLists.txt   -- SOURCE SPACE
      package.xml
    ...
    package_n/
      CATKIN_IGNORE    -- Optional empty file to exclude package_n from being processed
      CMakeLists.txt
      package.xml
    ...
  build/               -- BUILD SPACE
    CATKIN_IGNORE
  devel/              -- DEVELOPMENT SPACE (set by CATKIN_DEVEL_PREFIX)
    bin/
    etc/
    include/
    lib/
    share/
    .catkin
    env.bash
    setup.bash
    setup.sh
  ...
  install/            -- INSTALL SPACE (set by CMAKE_INSTALL_PREFIX)
    bin/
    etc/
    include/
    lib/
    share/
    .catkin
    env.bash
    setup.bash
    setup.sh
  ...
```

Figure 3.1: Structure of Catkin workspace

3.2.1 Source Space

The source space contains the source code of catkin packages. This is where you can extract/checkout/clone source code for the packages you want to build. Each folder within the source space contains one or more catkin packages. This space should remain unchanged by configuring, building, or installing. The root of the source space contains a symbolic link to catkin’s boiler-plate ‘toplevel’ CMakeLists.txt file. This file is invoked by CMake during the configuration of the catkin projects in the workspace. It can be created by calling `catkin_init_workspace` in the source space directory.

3.2.2 Build Space

The build space is where CMake is invoked to build the catkin packages in the source space. CMake and catkin keep their cache information and other intermediate files here. The build space does not have to be contained within the workspace nor does it have to be outside of the source space, but this is recommended.

3.2.3 Development (Devel) Space

The development space (or devel space) is where built targets are placed prior to being installed. The way targets are organized in the devel space is the same as their layout when they are installed. This provides a useful testing and development environment which does not require invoking the installation step. The location of the devel space is controlled by a catkin specific CMake variable called `CATKIN_DEVEL_PREFIX`, and it defaults to `<build space>/develspace`. This is the default behavior because it might be confusing to CMake users if they invoked `cmake ..` in a build folder and that modified things outside of the current directory. It is recommended, however, to set the devel space directory to be a peer of the build space directory.

3.2.4 Install Space

Once targets are built, they can be installed into the install space by invoking the `install` target, usually with `make install`. The install space does not have to be contained within the workspace. Since the install space is set by the `CMAKE_INSTALL_PREFIX`, it defaults to `/usr/local`, which you should not use.

3.2.5 Result space

When ever referring to a folder which can either be a development space or an install space the generic term result space is used.

3.3 Locosim

Locosim is a didactic framework to learn/test basic controllers schemes on quadruped robots and manipulators (UR5). It is composed by a roscontrol node called `ros_impedance_controller` (written in C++) that interfaces a python ROS node (where the controller is written) to a Gazebo simulator. For each controller, plotting / logging utilities are available to evaluate the results together with a configuration file (`LX_conf.py`) to change the controller parameters. [Github Reference](#)

3.4 Gazebo

Gazebo brings a fresh approach to simulation with a complete toolbox of development libraries and cloud services to make simulation easy. Iterate fast on your new physical designs in realistic environments with high fidelity sensors streams. Test control strategies in safety, and take advantage of simulation in continuous integration tests. It is composed by a set of open source development libraries, which encapsulate all the essentials, such as common math data types, logging, 3D mesh management, and asynchronous message passing. [reference](#)



Figure 3.2: Our simulation running on Gazebo

3.5 Rviz

Rviz is a virtualization environment that let us view what the robot is seeing, thinking and doing. Programming a robot can be very difficult without the possibility of debugging. There are two main ways to put data inside Rviz world (the combination of) ;

- Rviz understand sensors (like laser scans, cameras, corner frames)
- visualization markers (cubes, arrows, lines, colours, etc)

[reference](#)

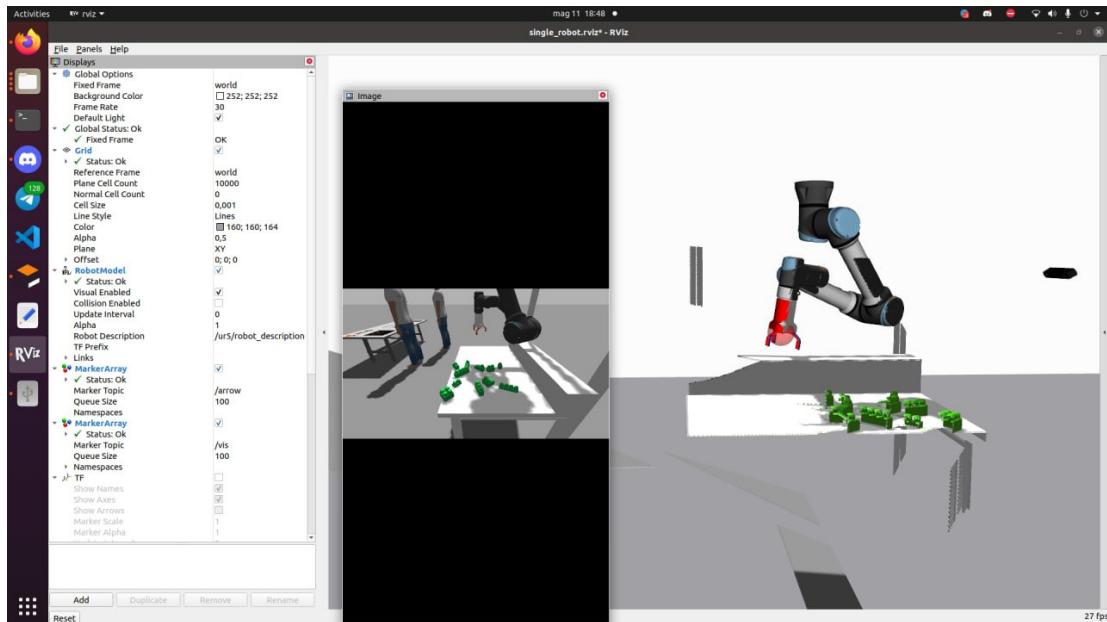


Figure 3.3: Our simulation running on Rviz

The black background window named "*image*" is the point of view of the ZED Camera, it's exactly what it is seeing.

4 Motion

The package motion is a Catkin package that contains 3 folders:

src It contains 2 executable files called *planner.cpp* and *movement.cpp*

msg It contains 2 message files called *legoFound.msg* and *legoTask.msg*

include It contains the *kinetics.h* library

4.1 Planner

Planner is the executable in charge of communicate with the *vision* package. The planner subscriber receives form this package through the vision publisher the lego class identified, its position and orientation. At the receiving side will arrive a message structured following the instructions inside the file stored in the motion package (*motion/msg/legoFound.msg*).

It has all the default destination coordinates for each class of lego.

AAAAAAAAAAAAAA INSERIRE SCREEN COORDINATE LEGO IN CODICE
PLANNER

4.1.1 Planner commands from vision

This commands are received from vision and informs how the planner module has to do. Then the planner computes and send specific commands to the movement module. That commands are explained in the movement section

No command: 0

aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa

Detect: 1

aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa

Quit: 2

aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa

4.1.2 motion/src/planner.cpp

4.1.3 motion/msg/legoFound.msg

<i>int32 command_id</i>	<i>float64 rot_roll</i>
<i>int32 lego_class</i>	<i>float64 rot_pitch</i>
<i>float64 coord_x</i>	<i>float64 rot_yaw</i>
<i>float64 coord_y</i>	<i>float64 date_time</i>
<i>float64 coord_z</i>	<i>string comment</i>
	<i>int32 send_ack</i>

```
void subDetectCommanderCallback(const motion::legoFound::ConstPtr &msg_detect); // reception and analysis of a received request
void pubDetectResulter(int risultato); // publish the result
void pubTaskCommander(bool s_ack); // publish the task
void subTaskResulterCallback(const motion::eventResult::ConstPtr &msg_event); // read ack sent by movement

Eigen::Vector3f camera2SimulationR(Eigen::Vector3f simul_lego_pos); // adapt camera point values in virtual world for robot
Eigen::Vector3f camera2RealR(Eigen::Vector3f camera_lego_pos); // adapt camera point values in real world for robot
void ungraspCommand(); // Enlarge robot fingers
void homingCommand(); // move the arm to default to avoid camera interferences

void catchCommand(Eigen::Vector3f position); // send the complete catch command to the moviment module

void selectClass(int lego_cl); // In relation to lego received from vision, load right lego parameters
double getTimeNow(); // Returns the current time
double getInterval(double start_t); // Returns the difference between the currentTime and the start time
```

Figure 4.1: Planner Functions Headers

By the ID class provided, it knows how to put the corresponding lego in its destination on the table. Once the planner has all the necessary info, exposes the legoTask.msg using the publisher to the [Movement](#) module.

4.2 Movement

It receives the *motion/msg/legoTask.msg* from the Planner and gets the info about:

- command to execute without caring about the class, most frequent is the 0: *catch_obj*
- initial coordinates
- diameter gripping and the final object position [optional]

Inside the *movement.cpp* there are basic instructions to move the UR5 robot referring to a library called *Kinetics*, (stored in /motion/include.kinetics.h).

Once the movement has terminated the task, send an ack to the planner, that checks that everything is done correctly and completely, then the planner communicates to the vision that requested the task the end of the jobs. At this point if the vision doesn't require further tasks, the workflow is complete, otherwise it will go on with the remaining requests.

4.2.1 motion/src/movement.cpp

```
// struct used to coordinate messages between planner and movement
struct ExecutingTask {
    int command_id;           // used to understand which task movement must execute
    double request_time;      // time to complete the task
    bool busy;                // It shows if the movement is performing a task or not
};
ExecutingTask planner_eseguendo;

Eigen::VectorXf joint_state_vector(6);
Eigen::VectorXf gripper_state_vector(3);

/*-----function headers-----*/

// reception and analysis of a received request. If it find an ack request in the end, it will publish eventResult message
void taskCommanderCallback(const motion::legoTask::ConstPtr &msg_taskCommand);
void pubTaskResulter(int risultato); // publish

// Returns the Trajectory vector using time and positions (start and end positions)
Eigen::Vector3f getTrajectory(double time, Eigen::Vector3f begin_position, Eigen::Vector3f final_position);
// it returns the joints speed
Eigen::VectorXf getJointSpeeds(Eigen::VectorXf joint_st, Eigen::VectorXf curr_position, Eigen::VectorXf destin_position, Eigen::VectorXf dest_v);

void updateJointStates(); // update joint state
float gripper2joints(float diameter); // returns the converted diameter of fingers
void nullCommandExecute(); // used when motion has no command to execute
void waitCommandExecute(int wait_time); // used to wait a command from planner
void moveProcedure(Eigen::Vector3f v_position, Eigen::Vector3f v_orientation, float dt); // used to move the arm
void graspObject(bool catchit); // used to grasp objects (open-close fingers)
void moveDefaultPosition(); // move the arm to a default position
void fastCatchProcedure(); // used when the command received is command_fast_catch: goes to the object in a fast way
void catchProcedure(); // used when the command received is command_catch: goes to the object using a full protocol
double getTimeNow(); // return the current time
double getInterval(double start_t); // return the interval between the start time and the current time
```

Figure 4.2: Movement Functions Headers

4.2.2 motion/msg/legoTask.msg

<i>int32 command_id</i>	<i>float64 rot_yaw</i>	<i>float64 dest_pitch</i>
<i>int32 real_robot</i>	<i>float64 gasp_diam</i>	<i>float64 dest_yaw</i>
<i>float64 coord_x</i>	<i>float64 dest_x</i>	<i>float64 ungasp_diam</i>
<i>float64 coord_y</i>	<i>float64 dest_y</i>	<i>int32 w_time</i>
<i>float64 coord_z</i>	<i>float64 dest_z</i>	<i>float64 date_time</i>
<i>float64 rot_roll</i>	<i>float64 dest_roll</i>	<i>string comment</i>
<i>float64 rot_pitch</i>		<i>int32 send_ack</i>

4.2.3 Movement commands from the planner module

This commands are received from the planner module and sent as execution to the robot arm.

No command: 0

aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa

Test: 1

aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa

Wait: 2

aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa

Move: 3

aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa

Grasp: 4

aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa

Ungrasp: 5

aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa

Default position: 6

aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa

Fast catch: 7

aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa

Catch: 8

aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa

4.3 Kinetics

Kinetics is the library that contains all the primitive instructions in order to calculate angles, spatial displacement, matrix and all the mathematical functions related.

Stored in `/motion/include/kinetics.h`, kinetics.h includes the necessary classes and functions for dense matrix operations, it provides a wide range of functionalities, including matrix algebraic operations.

```
// It calculates the jacobian Matrix related to the change of the reference system
Eigen::MatrixXf jacobMatrix(Eigen::VectorXf jo_ang) {

Eigen::VectorXf v_1(6);
v_1 << 0, -0.425, -0.3922, 0, 0, 0;

Eigen::VectorXf v_2(6);
v_2 << 0.1625, 0, 0, 0.1333, 0.0997, 0.0996 + 0.14;

Eigen::MatrixXf jac_mtx_1(6, 1);
jac_mtx_1 << v_2(4) * (cos(jo_ang(0)) * cos(jo_ang(4)) + cos(jo_ang(1)) * jo_ang(2) + jo_ang(3)) * sin(jo_ang(0)) * sin(jo_ang(4))) + v_2(2) * cos(jo
v_2(4) * (cos(jo_ang(4)) * sin(jo_ang(0)) - cos(jo_ang(1)) * jo_ang(2) + jo_ang(3)) * cos(jo_ang(0)) * sin(jo_ang(4))) + v_2(2) * sin(jo
0,
0,
0,
0,
1;

Eigen::MatrixXf jac_mtx_2(6, 1);
jac_mtx_2 << -cos(jo_ang(0)) * (v_1(2) * sin(jo_ang(1) + jo_ang(2)) + v_1(1) * sin(jo_ang(1)) + v_2(4) * (sin(jo_ang(1) + jo_ang(2)) * sin(jo_ang(3)
-sin(jo_ang(0)) * (v_1(2) * sin(jo_ang(1) + jo_ang(2)) + v_1(1) * sin(jo_ang(1)) + v_2(4) * (sin(jo_ang(1) + jo_ang(2)) * sin(jo_ang(3)
v_1(2) * cos(jo_ang(1) + jo_ang(2)) - (v_2(4) * sin(jo_ang(1) + jo_ang(2) + jo_ang(3) + jo_ang(4))) / 2 + v_1(1) * cos(jo_ang(1)) + (v_
sin(jo_ang(0)),
-cos(jo_ang(0)),
0;

Eigen::MatrixXf jac_mtx_3(6, 1);
jac_mtx_3 << cos(jo_ang(0)) * (v_2(4) * cos(jo_ang(1) + jo_ang(2) + jo_ang(3)) - v_1(2) * sin(jo_ang(1) + jo_ang(2)) + v_2(4) * sin(jo_ang(1) + jo_a
sin(jo_ang(0)) * (v_2(4) * cos(jo_ang(1) + jo_ang(2) + jo_ang(3)) - v_1(2) * sin(jo_ang(1) + jo_ang(2)) + v_2(4) * sin(jo_ang(1) + jo_a
v_1(2) * cos(jo_ang(1) + jo_ang(2)) - (v_2(4) * sin(jo_ang(1) + jo_ang(2) + jo_ang(3) + jo_ang(4))) / 2 + (v_2(4) * sin(jo_ang(1) + jo_a
-sin(jo_ang(0)),
-cos(jo_ang(0)),
0;
```

Figure 4.3: Kinetics.h

4.4 Acknowledgement system

It is available a feature that allows you to choose the operating mode. Since the communication between two different modules is done through messages with command IDs, each of these is defined in order to give back a variable called *send_ack*. If you want an execution results back to the caller, *send_ack* must be equal to **1**. It enforces the called module to complete the requested command and send back the results via event message. Viceversa if you just need the execution without caring about its completions, just leave "*send_ack*" equal to **0**.

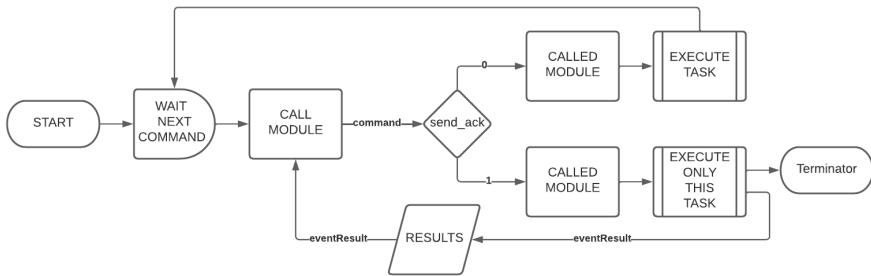


Figure 4.4: ACK Flow

4.5 spawnLego.cpp

Due to the need for thousands of images to get the best vision performance, we have developed a script that randomly generates the lego pieces on the table and takes pictures of the spawning area.

The script reads the parameters, spawns the legos, changes the colour of the legos that will be spawned, possibly removes a lego already generated and takes photos.

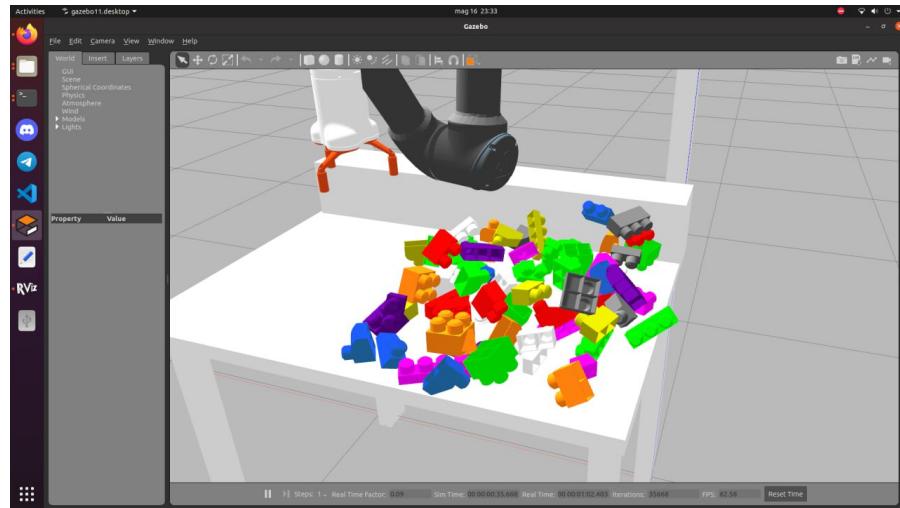


Figure 4.5: Spawn Lego Script

5 Vision

This module is in charge of practice the deep learning to create an “artificial neural network” that can learn and make intelligent decisions on its own, in our case the goal is to recognize the 11 different models (legos) starting from a dataset (pool of 1500 images).

5.1 Tool used

5.1.1 Yolov5

We have used YOLOv5, namely *You Only Look Once* as model for object detection.

YOLOv5 is a family of compound-scaled object detection models trained on the COCO dataset, and includes simple functionality for Test Time Augmentation (TTA), model ensembling, hyperparameter evolution, and export to ONNX, CoreML and TFLite.

It accepts URL, Filename, PIL, OpenCV, Numpy and PyTorch inputs, and returns detections in torch, pandas, and JSON output formats.

On June 25th, 2020, the first official version of YOLOv5 was released by [Ultralytics](#), a computer vision model used for detecting objects. [reference](#)

5.1.2 Pytorch

PyTorch is a fastest-growing Python machine learning framework based on the Torch library, used for applications such as computer vision and natural language processing, originally developed by Meta AI and now part of the Linux Foundation umbrella. It is free and open-source software released under the modified BSD license. Although the Python interface is more polished and the primary focus of development, PyTorch also has a C++ interface. [reference](#)

PyTorch provides two high-level features:

- Tensor computing with strong acceleration via graphics processing units (GPU)
- Deep neural networks built on a tape-based automatic differentiation system

5.1.3 Custom Roboflow version of Google Colab

In the official Ultralytics [Github Repository](#) there's a link that refers to [Google Colab](#) environment, where it is possible to use all the YOLOv5's tools in order to train the machine learning. We used the Roboflow version for some useful automations.

In this environment we had to upload the folder containing our images and values obtained using MakeSense or via Google Drive.

5.1.4 MakeSense AI

MakeSense AI is a free-to-use online tool for labeling photos. Thanks to the use of a browser it does not require any complicated installation - just visit the website and you are ready to go. It also doesn't matter which operating system you're running on - we do our best to be truly cross-platform. It is perfect for small computer vision deeplearning projects, making the process of preparing a dataset much easier and faster. Prepared labels can be downloaded in one of multiple supported formats. The application was written in TypeScript and is based on React/Redux duo. [reference](#)

Thanks to this tool we got the labels of all the dataset images.

The values of labels are obtained framing (selecting by hand a specific area of an image) and assigning to this frame a specific class, that is the name of the object family.

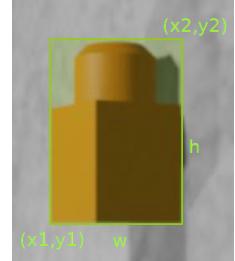
These values are stored into files called *annotations*, that are exportable in csv, xml or txt. The format depends on the platform that will be used for training (e.g. Roboflow uses .xml, instead yolov5 uses .txt)

obj class	X-position	Y-position	width	height
--------------	------------	------------	-------	--------

```
8 0.471170 0.624689 0.064544 0.105565
2 0.556368 0.553548 0.033563 0.064257
0 0.615318 0.621629 0.027539 0.065787
5 0.441910 0.558137 0.043890 0.070377
```

Figure 5.1: YOLOv5 labels values - Example

- For each images, a new file is created
- For each object, a new line is created
- For each different object in the image that we want to classify, we must define a new class, so in the example above we have 4 different objects selected.
- <object-class>: integer number of object { 0 - (classes-1) }
- <x> <y> <width> <height>: float values relative to width and height of image, it can be equal from (0.0 to 1.0)
 - For example: <x> = <absolute_x> / <image_width> or <height> = <absolute_height> / <image_height>
 - Attention: <x> <y> - are center of rectangle (are not top-left corner)



5.2 Steps

5.2.1 Create Dataset

First of all you have to create your dataset with a large pool of images (order of 1500/1700 images). The choices can't be made at random, you have to build a balanced dataset, it means that the number of images containing objects of class 1 must be similar to the number of images containing other classes. (No 1000 images of class 1 and 50 of class 3). Another important aspect is the variety of the data, that is in order to get the best results it's highly recommended to use images with different lights, shadows and spawn environment.

5.2.2 Dataset structure

If you want to build a reliable machine learning model, you need to split your dataset into the training, validation, and test sets. If you don't, your results will be biased, and you'll end up with a false impression of better model accuracy.

For training and testing purposes of our model, we should have our data broken down into three distinct dataset splits.

Training Set It's the set of data that is used to train and make the model learn the hidden features/patterns in the data. In each epoch, the same training data is fed to the neural network architecture repeatedly, and the model continues to learn the features of the data. The training set should have a diversified set of inputs so that the model is trained in all scenarios and can predict any unseen data sample that may appear hereafter.

Validation Set The validation set is a set of data, separate from the training set, that is used to validate our model performance during training. This validation process gives information that helps us tune the model's hyperparameters and configurations accordingly. It is like a critic telling us whether the training is moving in the right direction or not. The model is trained on the training set, and, simultaneously, the model evaluation is performed on the validation set after every epoch. The main idea of splitting the dataset into a validation set is to prevent our model from overfitting i.e., the model becomes really good at classifying the samples in the training set but cannot generalize and make accurate classifications on the data it has not seen before.

Test Set The test set is a separate set of data used to test the model after completing the training. It provides an unbiased final model performance metric in terms of accuracy, precision, etc. It answers the question of "How well does the model perform?"

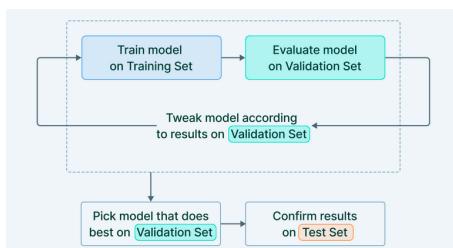


Figure 5.2: Train Vs. Validation Vs. Test

How to split the images

The creation of different samples and splits in the dataset helps us judge the true model performance. The dataset split ratio depends on the number of samples present in the dataset and the model. Some common inferences that can be derived on dataset split include:

- If there are several hyperparameters to tune, the machine learning model requires a larger validation set to optimize the model performance. Similarly, if the model has fewer or no hyperparameters, it would be easy to validate the model using a small set of data.
- If a model use case is such that a false prediction can drastically hamper the model performance—like falsely predicting cancer—it's better to validate the model after each epoch to make the model learn varied scenarios.
- With the increase in the dimension/features of the data, the hyperparameters of the neural network functions also increase making the model more complex. In these scenarios, a large split of data should be kept in training set with a validation set.

There is no optimal split percentage. One has to come to a split percentage that suits the requirements and meets the model's needs. However, there are two major concerns while deciding on the optimum split:

- If there is less training data, the machine learning model will show high variance in training.
- With less testing data/validation data, your model evaluation/model performance statistic will have greater variance.

Essentially, you need to come up with an optimum split that suits the need of the dataset/-model, but here's the rough standard split that you might encounter.

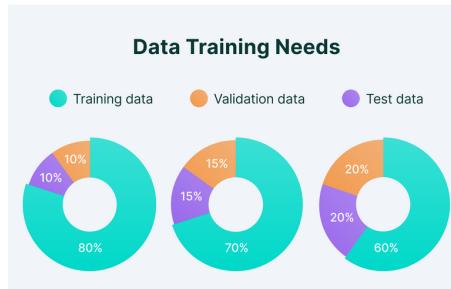


Figure 5.3: Recommended split

5.2.3 Create Labels for the annotations

We will use the [MakeSense AI](#) web tool in order to get the annotations. For each object that we want to detect we must associate to him a label, that will be then called class, since we have 11 types of legos: (an alternative is to load the *label.txt* file: one class per row)

0. X1-Y1-Z2
1. X1-Y2-Z1
2. X1-Y2-Z2
3. X1-Y2-Z2-CHAMFER
4. X1-Y2-Z2-TWINFILLET
5. X1-Y3-Z2
6. X1-Y3-Z2-FILLET
7. X1-Y4-Z1
8. X1-Y4-Z2
9. X2-Y2-Z2
10. X2-Y2-Z2-FILLET

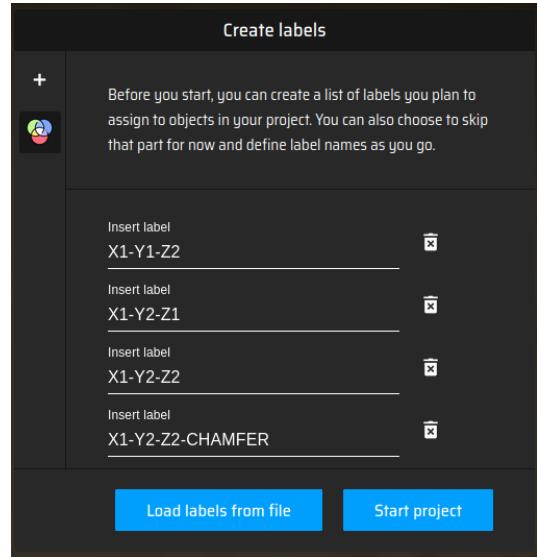


Figure 5.4: Labels in MakeSense AI

5.2.4 Create annotations

Once you have your data folders splitted in the way explained above, you can upload them to MakeSense AI in order to get the annotations. Keep in mind that you must be consistent in your choices (if a object is partial hidden, and the first time that occurs this situation you selected the hypothetic area, you must select in this way even for the next occurrences). Of course you must frame all the objects in a image, without forget anyone, so label every instance of every class.

5.2.5 Test and refine the dataset

Once we had all the annotations, we use the [Roboflow](#) tool in order to check our dataset health and get the best split. After the upload of the images Roboflow will check all the annotations, it will split the dataset into **train**, **valid** and **test**.

Roboflow suggests you tu use 640x046 format images, but it's very difficult to handle all the images in that way, but in the *generate* tab you are able to select all your preferences. A great way to operate is to follow the advices given by Roboflow. To improve the accuracy of the model there is the section *augmentation*, that performs tranforms of your existing images to create new variations and increase the number of images (same images flipped 90°, different brightness, etc).

5.2.6 Roboflow Dataset Settings

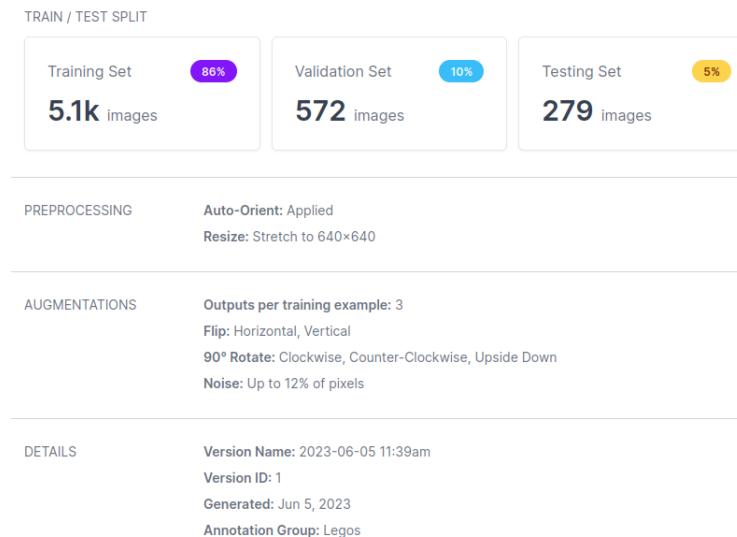


Figure 5.5: Dataset Settings

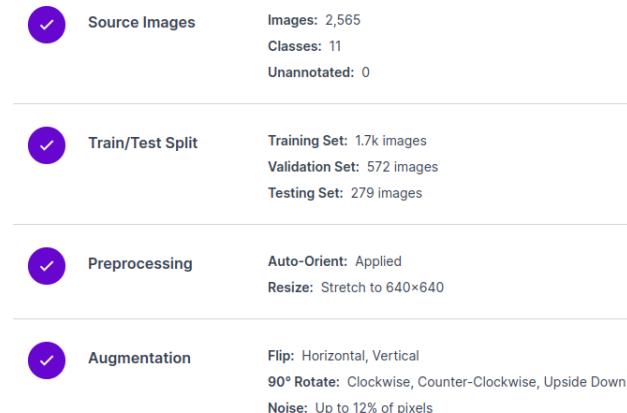


Figure 5.6: Additional settings

Now you are ready to launch the train machine, in the *versions* tab, wait for some uploads, then the train will start. Once the train is finished (300 epochs) you can try the model by uploading for instance a video. Furthermore they provide you the possibility to export the dataset elaborated through both a .zip file and a snippet code, this one is necessary to train our model into Google Colab.

5.2.7 Train into Google Colab provided by Ultralytics

Follow the [link](#) it will open a Google Colab environment customized by roboflow, in which you will use the previous exported snippet in order to train the model in Roboflow.

An other way to do it is by uploading the .zip dataset into the same Google Drive in which you have saved the Colab instance. (Note that you can have only one instance per account, furthermore the session will expire in 6 hours).

An alternative is to go in the official [Ultralytics Github Repository](#), in the *README.md*, section *Environment* you will find a link to the Google Colab environment.

Once signed-in, create a copy on drive of the Colab notebook, run the first part of code (setup section) in order to install the YOLOv5 folders, import the .zip and unzip it in the dataset folder.

Step 1 Clone the yolov5 repository and install all the necessary dependencies.

Step 2 Import dataset through the python code snippet previously copied in the roboflow platform, for instance our snippet is:

```
from roboflow import Roboflow
rf = Roboflow(api_key="1fFz7gZkuXvcnACwWiwR")
project = rf.workspace("roboticsiceunitn").project("robotics_ice23_unitn")
dataset = project.version(1).download("yolov5")
```

Step 3 Train the dataset with custom batch and epochs. We used 70 epochs and 64 batch, with *yolov5m.pt* weights.

Step 4 Test your model and then export /runs/ folder using a compressor (e.g. Zip) and save it in local because Colab will delete it.

```

+ Code + Text
Reconnect

[ ] Epoch GPU mem box loss obj loss cls loss Instances Size
64/69 7.396 0.01529 0.1245 0.01555 163 P R mAP50 mAP50-95: 100% 5/5 [00:04:06:00, 1.181it/s]
Class Images Instances P R
all 572 1420 0.963 0.999 0.971 0.872

Epoch GPU mem box loss obj loss cls loss Instances Size
65/69 7.396 0.01512 0.12173 0.01551 79 P R mAP50 mAP50-95: 100% 5/5 [00:04:06:00, 1.061it/s]
Class Images Instances P R
all 572 1420 0.956 0.916 0.972 0.871

Epoch GPU mem box loss obj loss cls loss Instances Size
66/69 7.396 0.01512 0.1219 0.01544 127 P R mAP50 mAP50-95: 100% 8/8 [00:04:06:00, 2.021it/s]
Class Images Instances P R
all 572 1420 0.956 0.923 0.973 0.875

Epoch GPU mem box loss obj loss cls loss Instances Size
67/69 7.396 0.01488 0.1185 0.01519 93 P R mAP50 mAP50-95: 100% 8/8 [00:04:06:00, 1.971it/s]
Class Images Instances P R
all 572 1420 0.968 0.914 0.973 0.877

Epoch GPU mem box loss obj loss cls loss Instances Size
68/69 7.396 0.01468 0.1185 0.01471 7 P R mAP50 mAP50-95: 100% 5/5 [00:04:06:00, 1.091it/s]
Class Images Instances P R
all 572 1420 0.968 0.914 0.973 0.877

Epoch GPU mem box loss obj loss cls loss Instances Size
69/69 7.396 0.01474 0.11933 0.01474 107 P R mAP50 mAP50-95: 100% 5/5 [00:05:06:00, 2.011it/s]
Class Images Instances P R
all 572 1420 0.962 0.916 0.974 0.88

70 epochs completed in 0.901 hours,
Optimized model saved to ./train/yolov5s/results/weights/best.pt, 14.9MB
Optimizer state dict saved to ./train/yolov5s/results/weights/best.pt, 14.99B

Validating run ./train/yolov5s/results/weights/best.pt...
Fusing layers...
custom YOLOv5s summary: 182 layers, 7723488 parameters, 0 gradients
    Images Instances P R mAP50 mAP50-95: 100% 5/5 [00:10:06:00, 2.18s/it]
X-Y-1-Z1 572 125 0.963 0.915 0.974 0.88
X1-Y1-Z2 572 99 0.981 0.919 0.98 0.875
X1-Y2-Z1 572 115 0.943 0.922 0.979 0.865
X1-Y2-Z2 572 144 0.943 0.922 0.979 0.865
X1-Y2-Z2-CHAMFER 572 125 0.965 0.882 0.96 0.86
X1-Y2-Z2-TWNINFILL 572 167 0.991 0.96 0.984 0.885
X1-Y2-Z2-TWNINFILL 572 172 0.991 0.961 0.984 0.885
X1-Y3-Z2-FILLET 572 125 0.98 0.872 0.973 0.871
X1-Y4-Z1 572 91 0.966 0.91 0.978 0.892
X1-Y4-Z2 572 148 0.963 0.913 0.983 0.893
X2-Y2-Z2-FILLET 572 149 0.934 0.947 0.985 0.895
X2-Y2-Z2-TWNINFILL 572 161 0.96 0.889 0.973 0.894

Results Saved To ./train/yolov5s/results
(CPU times: user 38.2 s, sys: 3.42 s, total: 33.7 s
Wall time: 55min 38s)

```

Figure 5.7: Colab results after training

5.3 Results

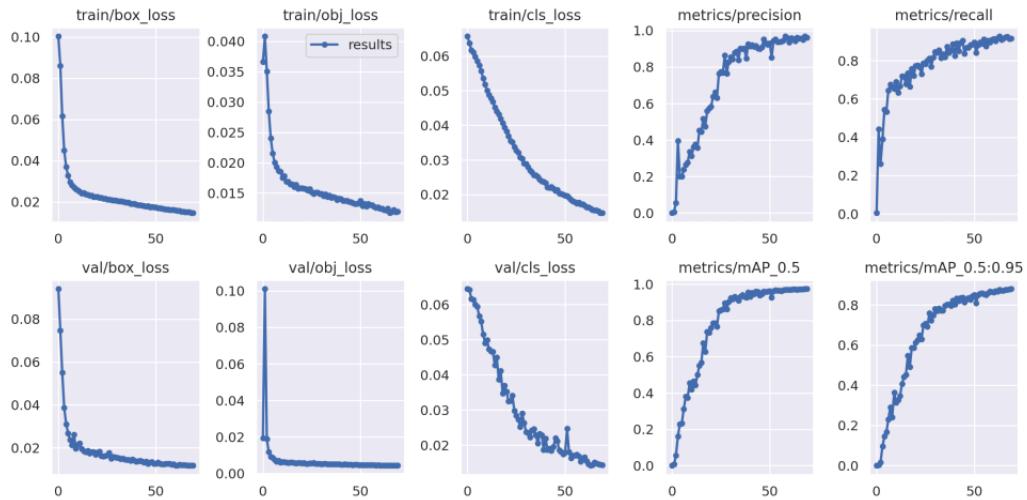


Figure 5.8: Results

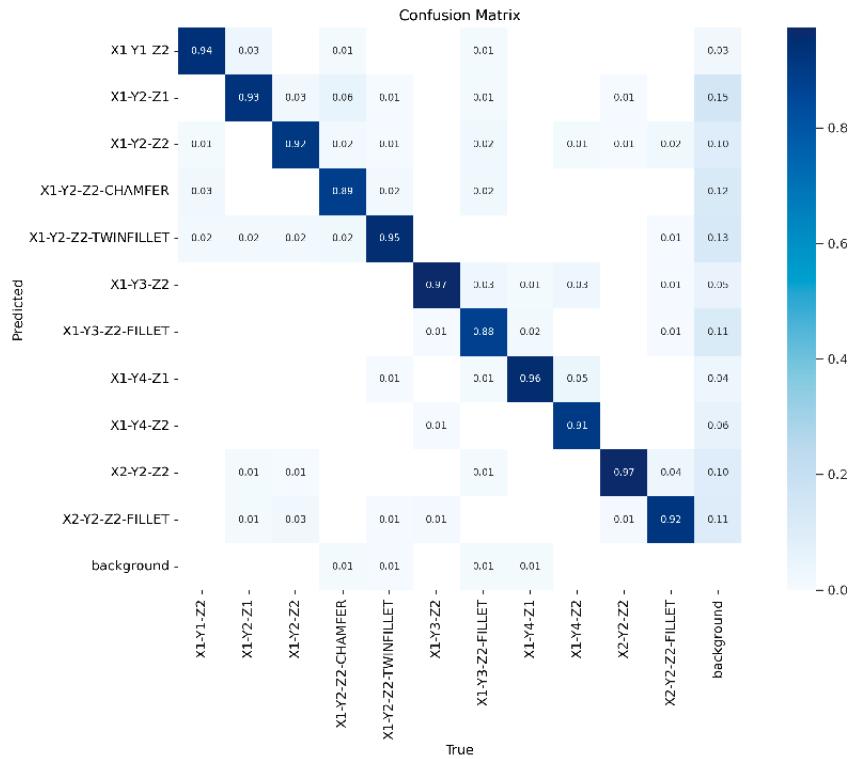


Figure 5.9: Confusion Matrix

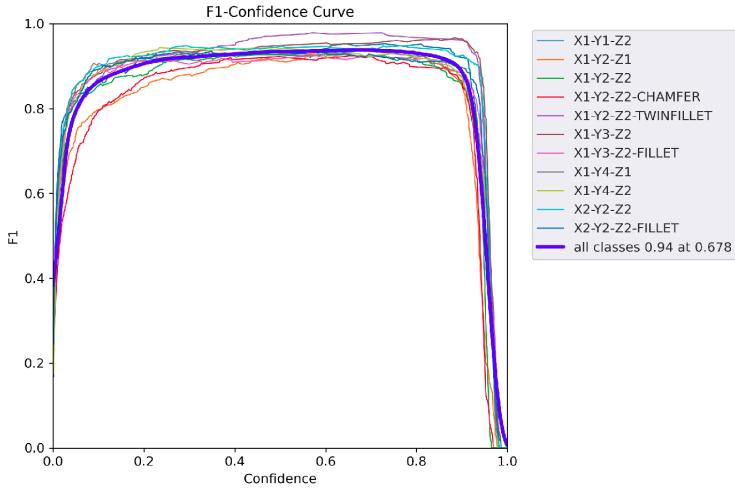


Figure 5.10: F1 Confidence Curve

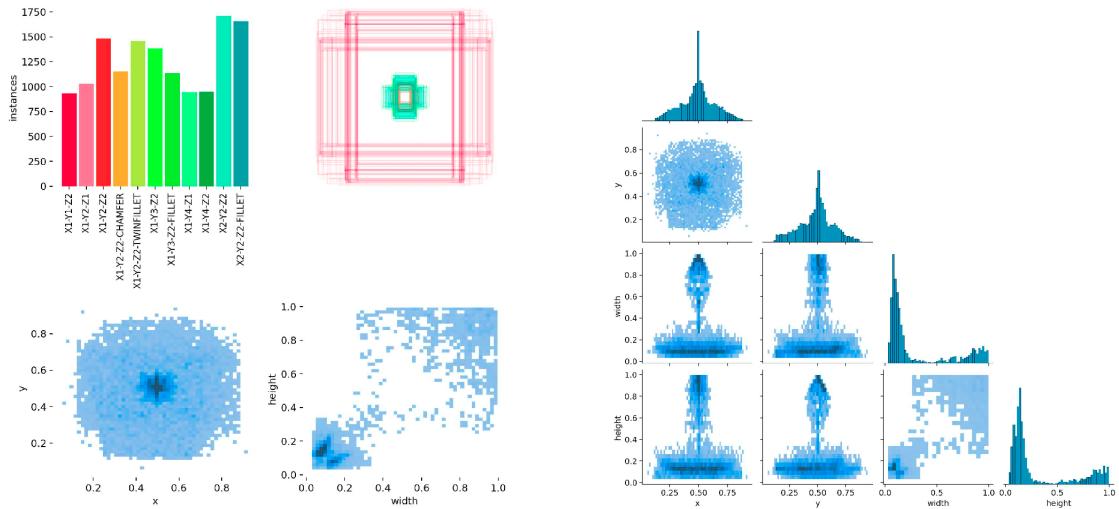


Figure 5.11: Labels & Labels Correlogram

5.4 Dataset

Our dataset in Roboflow available [here](#)

In the Colab platform, more in detail in the folder:

`/catkin_ws/src/vision/yolov5_22/runs/train/yolov5s_results/` you can find the results.

Obviously, due to the large size of our dataset, we won't include it inside the repository.

Note: we needed 22 attempts before get a good result. [Drive](#) with the previous weights.

5.5 In practice

5.5.1 Instance of assignment 2

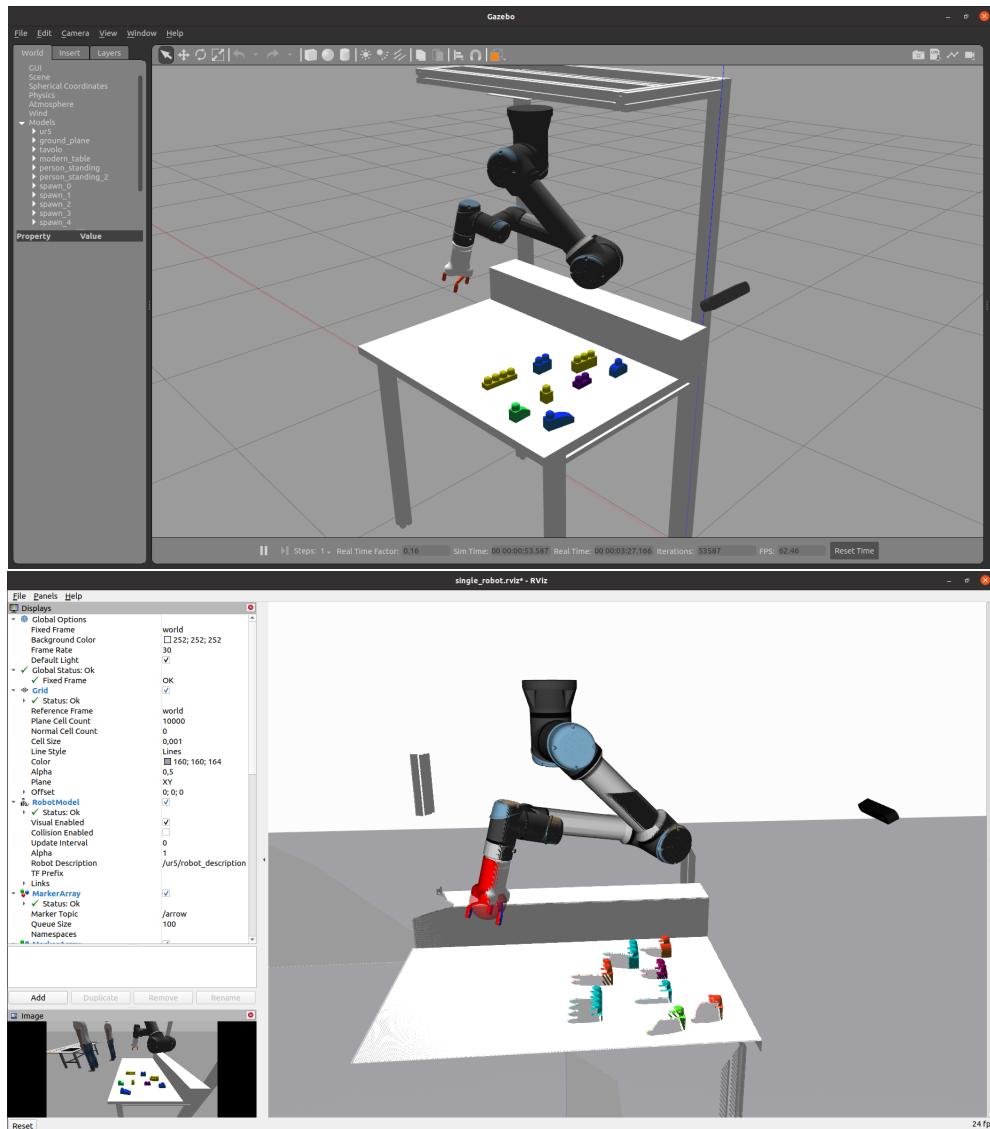


Figure 5.12: Instance of assignment 2

5.5.2 Normal Vision Results

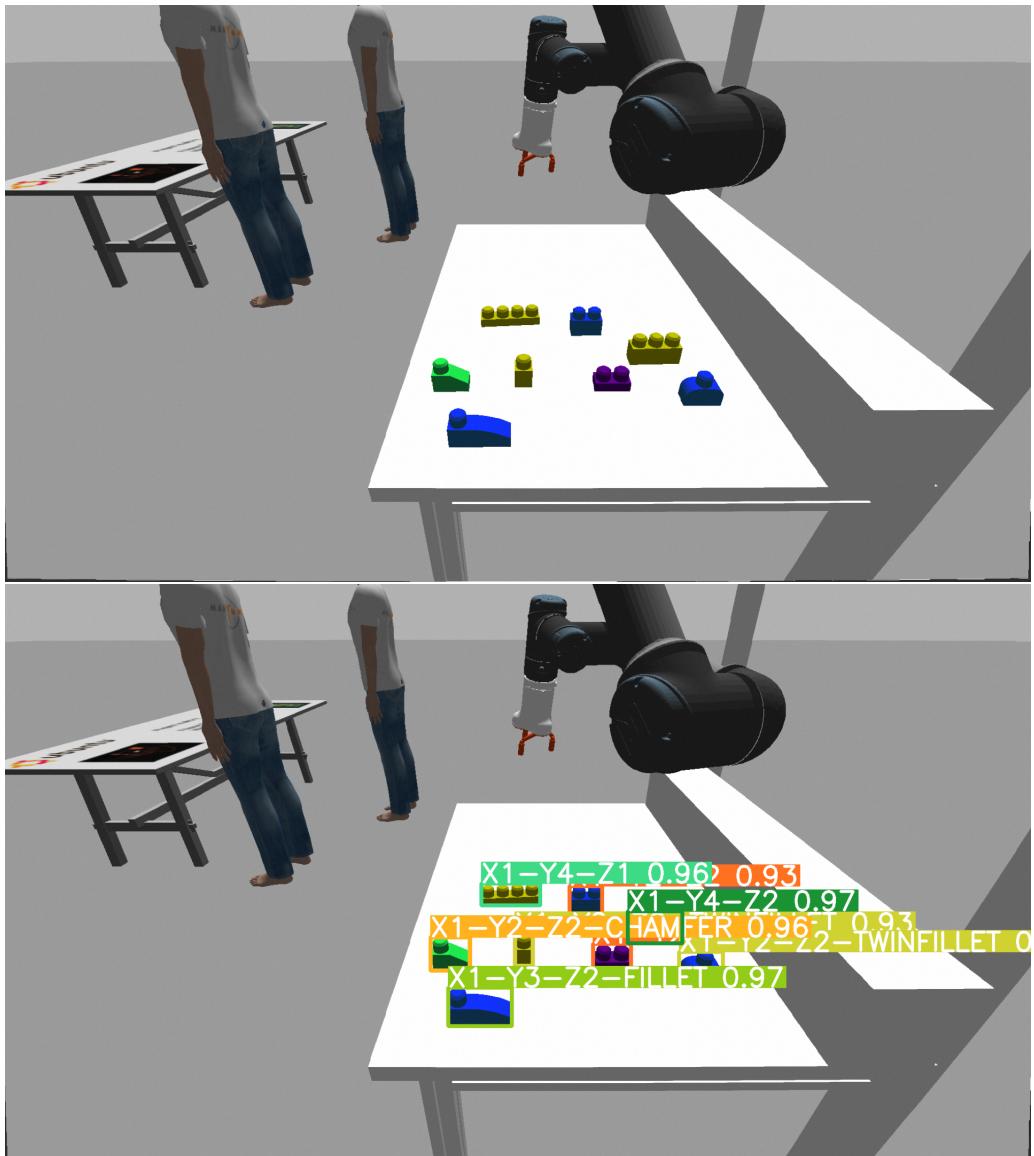


Figure 5.13: Normal Vision

5.5.3 ROI Vision Results

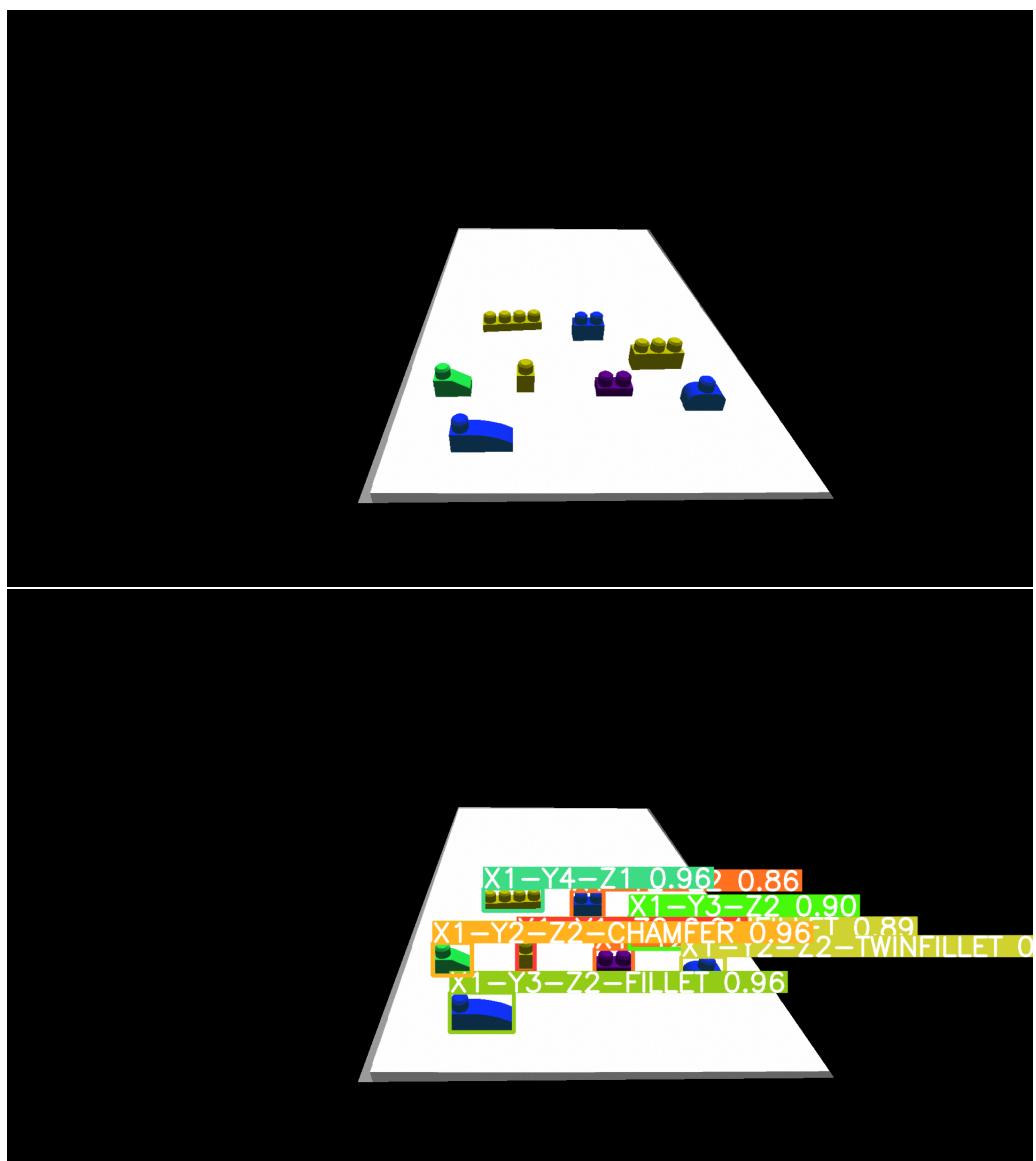


Figure 5.14: ROI Vision

6 Run The Project

After all the configuration you are able to run the project.

First of all move into catkin folder with **cd Robotics_ICE23_UNITN/catkin_ws/**.

Build the packages by typing **catkin_make install**.

If you get some issues, you just repeat the command.

Go back in the main folder **cd ..** and execute the start script **sh start.sh**.

Now you have to wait for Gazebo and Rviz to open. Once open, the UR5 robot homing procedure will start, which is shown in a dedicated terminal with red lettering.

Once you get the message "HOMING PROCEDURE ACCOMPLISHED," you can move to the other terminal open in the background and type any key to continue, then select the chosen assignment.